

CortexAI Platform Handbook

1. Overview

CortexAI is a full-stack generative AI platform. It combines a React and Vite frontend with a Node.js microservice backend covering chat, billing, authentication, and specialized AI agents. Requests are routed through a LangGraph supervisor graph that dispatches each user prompt to the correct agent.

2. Agent Catalog

The platform exposes eight agents: chat for general conversation, search for current-events lookups, coding for code generation and review, pdf for generating new documents, ppt for generating slide decks, image for image generation, vision for analyzing uploaded images, and pdf_rag for answering questions about an uploaded PDF.

3. Credit Costs

Each agent invocation consumes credits from the user's plan balance. The chat agent costs 1 credit per request. The search agent costs 5 credits per request. The coding agent costs 10 credits per request. The pdf agent costs 10 credits per request. The ppt agent costs 10 credits per request. The image agent costs 10 credits per request.

4. Rate Limits

To prevent abuse, each agent is limited on a rolling one-minute window. The chat agent allows 20 requests per minute. The search, coding, pdf, and ppt agents each allow 5 requests per minute. The image agent allows only 3 requests per minute, reflecting its higher compute cost.

5. Data Retention Policy

Files generated by the pdf, ppt, and image agents are stored in S3 and the download link expires after 24 hours. Uploaded PDFs used for retrieval-augmented question answering are embedded into a temporary vector collection and that collection is deleted immediately after the answer is generated, so no persistent index of an uploaded document is kept.

6. Support Policy

Standard support requests are answered within 2 business days. Enterprise customers on the Pro plan receive priority support with a response time of 4 hours. All support requests should be filed through the in-app help center rather than email.